

Tools for Debugging & Profiling

JSC OpenACC Course 2017

Contents

What you will learn. Hopefully.

- OpenACC can greatly speedup porting to GPU
- But many details hidden from user
- Compiler makes assumptions
- Programmer makes mistakes
- ⇒ Insight into program needed

Introduction

PGI Tools

Runtime Measurements

pgprof

NVIDIA Tools

cuda-memcheck

cuda-gdb

nvprof / pgprof

Visual Profiler

Tasks

Task 1

Task 2



```
$ ./spmv  
call to cuStreamSynchronize returned error 700: Illegal address  
during kernel execution
```

- Where does error come from?
- Is it an error at all?
- ... and how do I find out?

Important flags for PGI compiler

- Building for debugging
 - g Add debug information to executable; adds overhead → program performs slower
 - ta=tesla:lineinfo Add information to assembly to relate instructions to source code (*light debug info*)
- Check compiler output: -Minfo=accel

For NVIDIA's nvcc: -g, -lineinfo (, -G)

- Applications compiled with PGI compiler: Analyze via environment variables
- Maybe simplest/quickest check

PGI_ACC_TIME Lightweight profiler for time of data movement and kernels

PGI_ACC_NOTIFY Print information for GPU-related events.
Set to number, to print ...

- =1 ... kernel launches only
- =2 ... data transfers only
- =3 ... kernel launches and data transfers
- =4 ... region entry/exits only
- =5 ... region entry/exits and kernel launches
- =8 ... wait operations, synchronizations
- =16 ... (de)allocation of device memory

Usage: PGI_ACC_NOTIFY=3 ./app

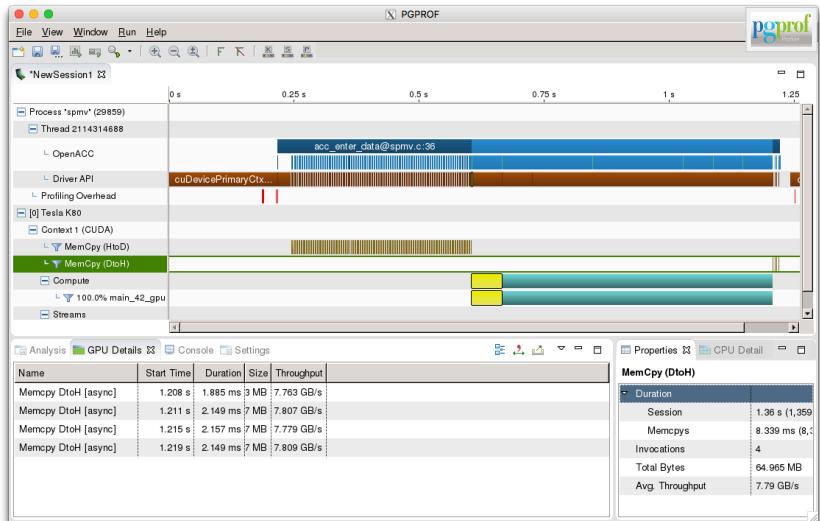
```
aherten@jrc0003:~/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/solution$ PGI_ACC_NOTIFY=3 ./spmv.bin
upload CUDA data  file=/homeb/zam/aherten/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/solution/spmv.c
  function=main line=36 device=0 threadid=1 variable=row_ptr bytes=15705192
upload CUDA data  file=/homeb/zam/aherten/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/solution/spmv.c
  function=main line=36 device=0 threadid=1 variable=row_ptr bytes=16777216
...
launch CUDA kernel file=/homeb/zam/aherten/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/solution/spmv.c
  function=main line=42 device=0 threadid=1 num_gangs=65535 num_workers=1 vector_length=128 grid=65535
  block=128 shared memory=1024
launch CUDA kernel file=/homeb/zam/aherten/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/solution/spmv.c
  function=main line=42 device=0 threadid=1 num_gangs=65535 num_workers=1 vector_length=128 grid=65535
  block=128 shared memory=1024
...
download CUDA data file=/homeb/zam/aherten/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/solution/spmv.c
  function=main line=59 device=0 threadid=1 variable=y bytes=14633352
download CUDA data file=/homeb/zam/aherten/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/solution/spmv.c
  function=main line=59 device=0 threadid=1 variable=y bytes=16777152
...
Runtime 0.172498 s.
```

- Graphical, interactive profiler
- Comes with PGI's compiler collection
- Nice visualizations, quick insight
- For OpenACC, OpenMP, CUDA
- Close to NVIDIA Visual Profiler

→ <https://www.pgroup.com/products/pgprof.htm>

PGPROF Graphical Performance Profiler

PGI's graphical profiler



- Memory error detector; similar to Valgrind's memcheck
- One of most helpful tools for error-finding
 - Out-of-bounds accesses
 - Kernels/API execution failures
 - Memory leaks
- Has sub-tools, via `cuda-memcheck --tool NAME`:
 - memcheck: Memory access checking (*default*)
 - racecheck: Shared memory hazard checking
 - Also: synccheck, initcheck
- Remember to compile program with debug information: `-g`

→ <http://docs.nvidia.com/cuda/cuda-memcheck/>

Start via cuda-memcheck app

```
aherten@jrl08:~/NVAL/Courses/OpenACC-2017/3-Debug/C/task2$ srun cuda-memcheck ./spmv.bin
===== Invalid __global__ read of size 8
=====   at 0x00000870 in /homeb/zam/aherten/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/spmv.c:50:main_4
=====   by thread (26,0,0) in block (123,0,0)
=====   Address 0x23aba45b60 is out of bounds
=====   Saved host backtrace up to driver entry point at kernel launch time
=====   Host Frame:/usr/lib64/nvidia/libcuda.so (cuLaunchKernel + 0x2c5) [0x204235]
=====   Host Frame:/usr/local/software/jureca/Stages/2017a/software/PGI/17.3-GCC-5.4.0/linux86-64/17.3
=====   Host Frame:/usr/local/software/jureca/Stages/2017a/software/PGI/17.3-GCC-5.4.0/linux86-64/17.3
=====   Host Frame:./spmv.bin [0x192b]
=====   Host Frame:/usr/lib64/libc.so.6 (__libc_start_main + 0xf5) [0x21b35]
=====   Host Frame:./spmv.bin [0xe69]
=====
===== Invalid __global__ read of size 8
=====   at 0x00000870 in /homeb/zam/aherten/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/spmv.c:50:main_4
=====   by thread (25,0,0) in block (123,0,0)
=====   Address 0x23aba45b58 is out of bounds
=====   Saved host backtrace up to driver entry point at kernel launch time
=====   Host Frame:/usr/lib64/nvidia/libcuda.so (cuLaunchKernel + 0x2c5) [0x204235]
=====   Host Frame:/usr/local/software/jureca/Stages/2017a/software/PGI/17.3-GCC-5.4.0/linux86-64/17.3
=====   Host Frame:/usr/local/software/jureca/Stages/2017a/software/PGI/17.3-GCC-5.4.0/linux86-64/17.3
=====   Host Frame:./spmv.bin [0x192b]
=====   Host Frame:/usr/lib64/libc.so.6 (__libc_start_main + 0xf5) [0x21b35]
```

- Powerful symbolic debugger for CUDA code
- Built on top of `gdb`
- Full usage: own course needed

```
cuda-gdb 101
```

```
run Starts application, give arguments with set args 1 2 ...
```

```
break L Create breakpoint
```

L: function name, line number LN, or FILE:LN

```
continue Continue running
```

```
print i Print content of i
```

```
info locals Print all currently set variables
```

```
info cuda threads Print current thread configuration
```

```
cuda thread N Switch context to thread number N
```

→ *cheat sheet*

→ <http://docs.nvidia.com/cuda/cuda-gdb/>

cuda-gdb can be used for OpenACC as well!

- Problem: Name of OpenACC-generated kernel?

→ Recipe: `strings ./app | grep .*_gpu | sort | uniq`

`strings ./app` Print occurrences of ≥ 4 printable characters

`grep .*_gpu` Search for `_gpu` line endings

`sort | uniq` Eliminate duplicates from list

- Examples of kernel names

Pattern: `function_line_gpu`

C `main_42_gpu`

Fortran `spmv_26_gpu`

Start via `cuda-gdb app → run`

Set breakpoint with `break func` or `break L` or `break file.c:L`

```
aherten@jrc0003:~/NVAL/Courses/OpenACC-2017/3-Debug/C/task2$ cuda-gdb spmv.bin
NVIDIA (R) CUDA Debugger
8.0 release
Portions Copyright (C) 2007-2016 NVIDIA Corporation
GNU gdb (GDB) 7.6.2
...
Reading symbols from /homeb/zam/aherten/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/spmv.bin...done.
(cuda-gdb) break main_42_gpu
Function "main_42_gpu" not defined.
Make breakpoint pending on future shared library load? (y or [n]) y
Breakpoint 1 (main_42_gpu) pending.
(cuda-gdb) run
Starting program: /homeb/zam/aherten/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/spmv.bin
...

[New Thread 0x2aab581ff700 (LWP 14126)]
[Switching focus to CUDA kernel 0, grid 1, block (0,0,0), thread (0,0,0), device 0, sm 12, warp 0, lane 0]

Breakpoint 1, main_42_gpu<<<(65535,1,1),(128,1,1)>>> (row_ptr=0x23051e0000, y=0x23a7400000, x=0x23a3600000,
43^I      for (int row=0; row<num_rows; ++row)
```

- Profiles CUDA kernels and API calls; also CPU code!
 - Suitable for OpenACC as well
 - pgprof: Very similar to nvprof, but different default options
 - Generate performance reports, timelines; measure events and metrics
- ⇒ Powerful complete tool for GPU application analysis
- <http://docs.nvidia.com/cuda/profiler-users-guide/>

Start via nvprof ./app

```

aherten@jrc0003:~/NVAL/Courses/OpenACC-2017/3-Debug/C/task2/solution$ nvprof ./spmv.bin
==18860== NVPROF is profiling process 18860, command: ./spmv.bin
Runtime 0.250184 s.
==18860== Profiling application: ./spmv.bin
==18860== Profiling result:
Time(%)    Time      Calls      Avg      Min      Max  Name
86.96%    2.50132s      10    250.13ms  250.09ms  250.18ms  main_42_gpu
12.74%    366.33ms     168    2.1805ms  234.17us  2.3417ms  [CUDA memcpy HtoD]
0.30%     8.6459ms      4     2.1615ms  1.9484ms  2.2338ms  [CUDA memcpy DtoH]

==18860== API calls:
Time(%)    Time      Calls      Avg      Min      Max  Name
81.21%    2.50627s     12    208.86ms  2.2347ms  250.18ms  cuStreamSynchronize
9.33%     287.85ms      2    143.93ms    272ns    287.85ms  cuDevicePrimaryCtxRetain
5.51%     170.02ms      1     170.02ms  170.02ms  170.02ms  cuDevicePrimaryCtxRelease
2.31%     71.431ms     172     415.29us  2.4300us  2.2410ms  cuEventSynchronize
0.80%     24.810ms      1     24.810ms  24.810ms  24.810ms  cuMemHostAlloc
0.41%     12.720ms      6     2.1199ms  530.76us  3.2406ms  cuMemAlloc
0.34%     10.440ms      1     10.440ms  10.440ms  10.440ms  cuMemFreeHost
0.04%      1.1568ms     168     6.8850us  5.8200us  29.560us  cuMemcpyHtoDAsync
0.02%      609.58us      1     609.58us  609.58us  609.58us  cuMemAllocHost
...

```

Start via pgprof ./app

```

aherten@jrc0003:~/NVAL/Courses/OpenACC-2017/3-Debug/Fortran/task2/solution$ pgprof ./spmv.bin
==9116== PGPROF is profiling process 9116, command: ./spmv.bin
Runtime: 0.624193 s.
==9116== Profiling application: ./spmv.bin
==9116== Profiling result:
Time(%)   Time      Calls      Avg      Min      Max   Name
94.30%   6.24126s      10   624.13ms  624.05ms  624.20ms  spmv_26_gpu
 5.57%   368.82ms     173   2.1319ms  4.7680us  2.3071ms  [CUDA memcpy HtoD]
 0.13%    8.6119ms      4    2.1530ms  1.9330ms  2.2314ms  [CUDA memcpy DtoH]

==9116== API calls:
Time(%)   Time      Calls      Avg      Min      Max   Name
90.96%   6.24617s      12   520.51ms  2.2297ms  624.21ms  cuStreamSynchronize
...

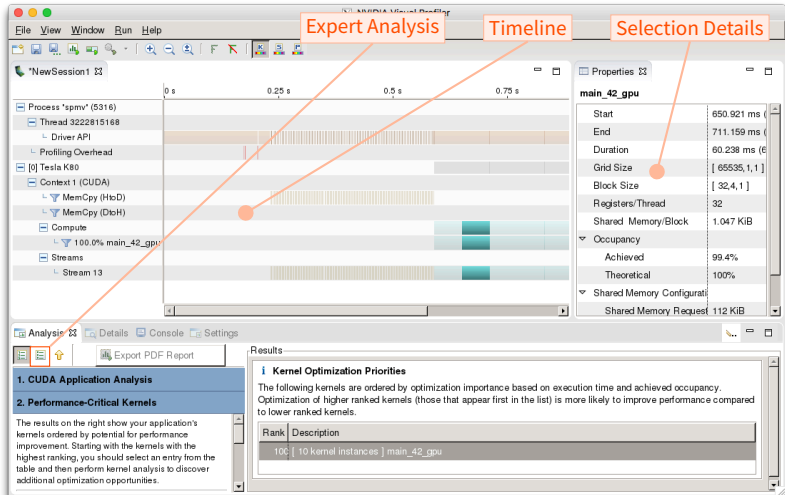
==9116== OpenACC (excl):
Time(%)   Time      Calls      Avg      Min      Max   Name
93.66%   6.24139s      10   624.14ms  624.06ms  624.22ms  acc_wait@spmv.F03:26
 6.03%   401.76ms      1    401.76ms  401.76ms  401.76ms  acc_enter_data@spmv.F03:22
 0.17%   11.199ms      1    11.199ms  11.199ms  11.199ms  acc_exit_data@spmv.F03:22
 0.06%    3.9573ms      1    3.9573ms  3.9573ms  3.9573ms  acc_wait@spmv.F03:41
...

```


- Timeline view of all things GPU (API calls, kernels, memory)
→ study stages and interplay of application
- View launch and run configurations
- Guided and unguided analysis, with (among others):
 - Performance limiters
 - Kernel and execution properties
 - Memory access patterns
- NVIDIA version (nvvp) or PGI version (pgprof without arguments)

→ <https://developer.nvidia.com/nvidia-visual-profiler>

Start via nvvp → *File* ↪ *New Session*



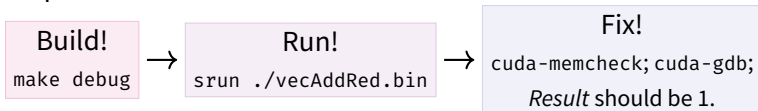
Task 1

TASK

C

FORTRAN

- Location of tasks:
Debugging/tasks/{C,Fortran}/task1/
- Vector addition and reduction: $\vec{c} = \vec{a} + \vec{b} \rightarrow \gamma = \sum_i c_i$
- Steps




JURECA Getting Started

```
module load PGI CUDA
salloc --reservation=oacc17 --partition=gpus --nodes=1 --time=1:30:00
↪ --gres=mem128,gpu:4
srun cuda-memcheck ./vecAddRed.bin
```


- Sparse Matrix-Vector Product ($SpMV$): $\vec{x} = A\vec{y}$
- CSR data layout

	0	1	2	3	4
0	-2	1	0	0	0
1	1	-2	1	0	0
2	0	1	-2	1	0
3	0	0	1	-2	1
4	0	0	0	1	-2

	0	1	2	3	4
0	-2	1			
1	1	-2	1		
2		1	-2	1	
3			1	-2	1
4				1	-2



- Sparse Matrix-Vector Product ($SpMV$): $\vec{y} = \mathbf{A}\vec{x}$
- CSR data layout



JURECA Getting Started

```
module load PGI CUDA
salloc --reservation=oacc17 --partition=gpus --nodes=1 --time=1:30:00
↪ --gres=mem128,gpu:4
srun cuda-memcheck ./spmv.bin
```

- All the CUDA debugging and performance measurement tools work
 - pgprof
 - cuda-memcheck
 - cuda-gdb
 - nvprof
 - Visual Profiler
- Sometimes, a little digging is needed to find automatically-generated function names

Happy Debugging!
a.herten@fz-juelich.de

Appendix

Glossary

CUDA Computing platform for **GPUs** from NVIDIA. Provides, among others, CUDA C/C++. 7, 11, 14, 22

NVIDIA US technology company creating **GPUs**. 4, 7, 17, 24

OpenACC Directive-based programming, primarily for many-core machines. 1, 2, 7, 12, 14

OpenMP Directive-based programming, primarily for multi-threaded machines. 7

PGI Compiler creators. Formerly *The Portland Group, Inc.*; since 2013 part of **NVIDIA**. 4, 5, 6, 7, 8, 17